

Optimizing C++'s



Emergent Structures



Emergent Structures

Compilers

Principles, Techniques,
and Tools





Compiler

Frontend

Optimizer

Code Generator

Optimization 101

```
extern "C" int atoi(const char *);  
int main(int argc, char **argv) {  
    if (argc != 3)  
        return -1;  
  
    return atoi(argv[1]) + atoi(argv[2]);  
}
```

```

TranslationUnitDecl 0x5332240 <<invalid sloc>>
| -LinkageSpecDecl 0x5332bf0 <hello_world.cpp:1:1, col:33> C
|   \-FunctionDecl 0x5332d40 <col:12, col:33> atoi 'int (const char *)'
|     \-ParmVarDecl 0x5332c80 <col:21, col:33> 'const char *'
| -FunctionDecl 0x535e680 <line:2:1, line:7:1> main 'int (int, char **)'
|   | -ParmVarDecl 0x5332e00 <line:2:10, col:14> argc 'int'
|   | -ParmVarDecl 0x5332ed0 <col:20, col:27> argv 'char **'
|   \-CompoundStmt 0x535ebb8 <col:33, line:7:1>
|     | -IfStmt 0x535e818 <line:3:3, line:4:13>
|     |   | -<<<NULL>>>
|     |   | -BinaryOperator 0x535e790 <line:3:7, col:15> '_Bool' '!='
|     |   |   | -ImplicitCastExpr 0x535e778 <col:7> 'int' <LValueToRValue>
|     |   |   |   \-DeclRefExpr 0x535e730 <col:7> 'int' lvalue ParmVar 0x5332e00 'argc' 'int'
|     |   |   |   \-IntegerLiteral 0x535e758 <col:15> 'int' 3
|     |   | -ReturnStmt 0x535e7f8 <line:4:5, col:13>
|     |   |   \-UnaryOperator 0x535e7d8 <col:12, col:13> 'int' prefix '-'
|     |   |     \-IntegerLiteral 0x535e7b8 <col:13> 'int' 1
|     |   \-<<<NULL>>>
|     \-ReturnStmt 0x535eb98 <line:6:3, col:38>
|       \-BinaryOperator 0x535eb70 <col:10, col:38> 'int' '+'
|         | -CallExpr 0x535e990 <col:10, col:22> 'int'
|         |   | -ImplicitCastExpr 0x535e978 <col:10> 'int (*)(const char *)' <FunctionToPointerDecay>
|         |   |   \-DeclRefExpr 0x535e928 <col:10> 'int (const char *)' lvalue Function 0x5332d40 'atoi' '...'
|         |   \-ImplicitCastExpr 0x535e9d8 <col:15, col:21> 'const char *' <NoOp>
|         |     \-ImplicitCastExpr 0x535e9c0 <col:15, col:21> 'char *' <LValueToRValue>
|         |       \-ArraySubscriptExpr 0x535e900 <col:15, col:21> 'char *' lvalue
|         |         | -ImplicitCastExpr 0x535e8e8 <col:15> 'char **' <LValueToRValue>
|         |         |   \-DeclRefExpr 0x535e8a0 <col:15> 'char **' lvalue ParmVar 0x5332ed0 'argv' 'char **'
|         |         \-IntegerLiteral 0x535e8c8 <col:20> 'int' 1
|         \-CallExpr 0x535eb10 <col:26, col:38> 'int'
|           | -ImplicitCastExpr 0x535eaf8 <col:26> 'int (*)(const char *)' <FunctionToPointerDecay>
|           |   \-DeclRefExpr 0x535ead0 <col:26> 'int (const char *)' lvalue Function 0x5332d40 'atoi' 'int (const

```

```

define i32 @main(i32 %argc, i8** %argv) {
entry:
    %retval = alloca i32, align 4
    %argc.addr = alloca i32, align 4
    %argv.addr = alloca i8**, align 8
    store i32 0, i32* %retval
    store i32 %argc, i32* %argc.addr, align 4
    store i8** %argv, i8*** %argv.addr, align 8
    %0 = load i32* %argc.addr, align 4
    %cmp = icmp ne i32 %0, 3
    br i1 %cmp, label %if.then, label %if.end

if.then:                                     ; preds = %entry
    store i32 -1, i32* %retval
    br label %return

if.end:                                     ; preds = %entry
    %1 = load i8*** %argv.addr, align 8
    %arrayidx = getelementptr inbounds i8** %1, i64 1
    %2 = load i8** %arrayidx, align 8
    %call = call i32 @atoi(i8* %2)
    %3 = load i8*** %argv.addr, align 8
    %arrayidx1 = getelementptr inbounds i8** %3, i64 2
    %4 = load i8** %arrayidx1, align 8
    %call2 = call i32 @atoi(i8* %4)
    %add = add nsw i32 %call, %call2
    store i32 %add, i32* %retval
    br label %return

return:                                     ; preds = %if.end, %if.then
    %5 = load i32* %retval
    ret i32 %5
}

```

A brief digression to describe
LLVM's IR...

```
declare i32 @g(i32 %x)

define i32 @f(i32 %a, i32 %b) {
entry:
    %c = add i32 %a, %b
    %d = call i32 @g(i32 %c)
    %e = add i32 %c, %d
    ret i32 %e
}
```

```
declare i32 @g(i32 %x)

define i32 @f(i32 %a, i32 %b, i1 %flag) {
entry:
    %c = add i32 %a, %b
    br i1 %flag, label %then, label %else

then:
    %d = call i32 @g(i32 %c)
    ret i32 %d

else:
    ret i32 %c
}
```

```
declare i32 @g(i32 %x)

define i32 @f(i32 %a, i32 %b, i1 %flag) {
entry:
    %c = add i32 %a, %b
    br i1 %flag, label %then, label %end

then:
    %d = call i32 @g(i32 %c)
    br label %end

end:
    %result = phi i32 [ %entry, %c ],
                  [ %then, %d ]
    ret i32 %result
}
```

Ok, where were we...

IR for hello world

```

define i32 @main(i32 %argc, i8** %argv) {
entry:
    %retval = alloca i32, align 4
    %argc.addr = alloca i32, align 4
    %argv.addr = alloca i8**, align 8
    store i32 0, i32* %retval
    store i32 %argc, i32* %argc.addr, align 4
    store i8** %argv, i8*** %argv.addr, align 8
    %0 = load i32* %argc.addr, align 4
    %cmp = icmp ne i32 %0, 3
    br i1 %cmp, label %if.then, label %if.end

if.then:                                     ; preds = %entry
    store i32 -1, i32* %retval
    br label %return

if.end:                                     ; preds = %entry
    %1 = load i8*** %argv.addr, align 8
    %arrayidx = getelementptr inbounds i8** %1, i64 1
    %2 = load i8** %arrayidx, align 8
    %call = call i32 @atoi(i8* %2)
    %3 = load i8*** %argv.addr, align 8
    %arrayidx1 = getelementptr inbounds i8** %3, i64 2
    %4 = load i8** %arrayidx1, align 8
    %call2 = call i32 @atoi(i8* %4)
    %add = add nsw i32 %call, %call2
    store i32 %add, i32* %retval
    br label %return

return:                                     ; preds = %if.end, %if.then
    %5 = load i32* %retval
    ret i32 %5
}

```

Optimization does more than just
make code faster...

Step 1: Cleanup

```

extern "C" int atoi(const char *);
int main(int argc, char **argv) {
    if (argc != 3)
        return -1;

    return atoi(argv[1]) +
           atoi(argv[2]);
}

```

```

define i32 @main(i32 %argc, i8** %argv) {
entry:
    %retval = alloca i32, align 4
    %argc.addr = alloca i32, align 4
    %argv.addr = alloca i8**, align 8
    store i32 0, i32* %retval
    store i32 %argc, i32* %argc.addr, align 4
    store i8** %argv, i8*** %argv.addr, align 8
    %0 = load i32* %argc.addr, align 4
    %cmp = icmp ne i32 %0, 3
    br i1 %cmp, label %if.then, label %if.end

if.then:
    store i32 -1, i32* %retval
    br label %return

if.end:
    %1 = load i8*** %argv.addr, align 8
    %arrayidx = getelementptr i8** %1, i64 1
    %2 = load i8** %arrayidx, align 8
    %call = call i32 @atoi(i8* %2)
    %3 = load i8*** %argv.addr, align 8
    %arrayidx1 = getelementptr i8** %3, i64 2
    %4 = load i8** %arrayidx1, align 8
    %call2 = call i32 @atoi(i8* %4)
    %add = add nsw i32 %call, %call2
    store i32 %add, i32* %retval
    br label %return

return:
    %5 = load i32* %retval
    ret i32 %5
}

```

```

extern "C" int atoi(const char *);
int main(int argc, char **argv) {
    if (argc != 3)
        return -1;

    return atoi(argv[1]) +
        atoi(argv[2]);
}

```

```

define i32 @main(i32 %argc, i8** %argv) {
entry:
    %cmp = icmp ne i32 %argc, 3
    br i1 %cmp, label %if.then, label %if.end

if.then:
    br label %return

if.end:
    %arrayidx = getelementptr i8** %argv, i64 1
    %arrayval = load i8** %arrayidx, align 8
    %call = call i32 @atoi(i8* %arrayval)
    %arrayidx1 = getelementptr i8** %argv, i64 2
    %arrayval1 = load i8** %arrayidx1, align 8
    %call12 = call i32 @atoi(i8* %arrayval1)
    %add = add nsw i32 %call, %call12
    br label %return

return:
    %retval.0 = phi i32 [ -1, %if.then ],
                    [ %add, %if.end ]
    ret i32 %retval.0
}

```

Step 2: Canonicalization

```
int x = y;  
if (!flag)  
    x = z;
```

```
int x;  
if (flag)  
    x = y;  
else  
    x = z;
```

```
if (flag)  
    z = y;  
int x = z;
```

```
int x = flag? y : z;
```

```

extern "C" int atoi(const char *);
int main(int argc, char **argv) {
    if (argc != 3)
        return -1;

    return atoi(argv[1]) +
        atoi(argv[2]);
}

```

```

define i32 @main(i32 %argc, i8** %argv) {
entry:
    %cmp = icmp ne i32 %argc, 3
    br i1 %cmp, label %if.then, label %if.end

if.then:
    br label %return

if.end:
    %arrayidx = getelementptr i8** %argv, i64 1
    %arrayval = load i8** %arrayidx, align 8
    %call = call i32 @atoi(i8* %arrayval)
    %arrayidx1 = getelementptr i8** %argv, i64 2
    %arrayval1 = load i8** %arrayidx1, align 8
    %call2 = call i32 @atoi(i8* %arrayval1)
    %add = add nsw i32 %call, %call2
    br label %return

return:
    %retval.0 = phi i32 [ -1, %if.then ],
                    [ %add, %if.end ]
    ret i32 %retval.0
}

```

```

extern "C" int atoi(const char *);
int main(int argc, char **argv) {
    if (argc != 3)
        return -1;

    return atoi(argv[1]) +
        atoi(argv[2]);
}

```

```

define i32 @main(i32 %argc, i8** %argv) {
entry:
    %cmp = icmp eq i32 %argc, 3
    br i1 %cmp, label %if.end, label %return

if.end:
    %arrayidx = getelementptr i8** %argv, i64 1
    %arrayval = load i8** %arrayidx, align 8
    %call = call i32 @atoi(i8* %arrayval)
    %arrayidx1 = getelementptr i8** %argv, i64 2
    %arrayval1 = load i8** %arrayidx1, align 8
    %call2 = call i32 @atoi(i8* %arrayval1)
    %add = add nsw i32 %call2, %call
    br label %return

return:
    %retval.0 = phi i32 [ %add, %if.end ],
        [ -1, %entry ]
    ret i32 %retval.0
}

```

Step 3: Collapse Abstractions

Three key abstractions:


1. Functions, calls, and the call graph.
2. Memory, loads, and stores.
3. Loops.

But no one really cares about
loops in C++... Right?

Fortran

Because


if you really care...




Also, we know how to optimize
loops


What about the other
fundamentals?


Let's look at collapsing function
calls






Bottom-up SCC-based call graph
walk







How does the optimizer evaluate complexity?

```
int g(double x, double y, double z);
```

```
int f(struct S* s, double y, double x) {  
    return g(x, y, s->z);  
}
```

```
void fancy_sort(vector<int> &v) {  
    if (v.size() <= 1)  
        return;  
    if (v.size() == 2) {  
        if (v.front() >= v.back())  
            swap(v.front(), v.back());  
        return;  
    }  
    std::sort(v.begin(), v.end());  
}
```

This doesn't always work though!

```
int hash(hash_state &h) {  
    // Some complex code on 'h'  
    return /* final value */;  
}
```

```
template <typename T, typename ...Ts>  
int hash(hash_state &h, T arg, Ts ...args) {  
    // Complex code to put 'arg'  
    // into the 'h' state...  
    return hash(h, args...);  
}
```

Let's look at memory, loads, and
stores...


```

%S = type { i32, i32, i32 }
declare i32 @g(i32 %x)
define i32 @f(i32 %a, i32 %b, i1 %flag) {
entry:
    %mem = alloca %S
    %c = add i32 %a, %b
    %addr0 = getelementptr %S* %mem,
                i32 0, i32 0
    store i32 %c, i32* %addr0
    %addr1 = getelementptr %S* %mem,
                i32 0, i32 1
    store i32 0, i32* %addr1
    %addr2 = getelementptr %S* %mem,
                i32 0, i32 2
    store i32 0, i32* %addr2
    br i1 %flag, label %then, label %end

then:
    %d = call i32 @g(i32 %c)
    store i32 %d, i32* %addr1
    %e = call i32 @g(i32 %a)
    store i32 %e, i32* %addr2
    br label %end

end:
    %val0 = load i32* %addr0
    %val1 = load i32* %addr1
    %val2 = load i32* %addr2
    %f = add i32 %val0, %val1
    %result = add i32 %f, %val2
    ret i32 %result
}

```

```

declare i32 @g(i32)

define i32 @f(i32 %a, i32 %b, i1 %flag) {
entry:
    %c = add i32 %a, %b
    br i1 %flag, label %then, label %end

then:
    %d = call i32 @g(i32 %c)
    %e = call i32 @g(i32 %a)
    br label %end

end:
    %mem.sroa.1.0 =
        phi i32 [ %d, %then ], [ 0, %entry ]
    %mem.sroa.2.0 =
        phi i32 [ %e, %then ], [ 0, %entry ]
    %f = add i32 %c, %mem.sroa.1.0
    %result = add i32 %f, %mem.sroa.2.0
    ret i32 %result
}

```

Not all memory is non-escaping
stack addresses... But that's OK

What happens when these
abstractions are *combined*?

```
int f(int a, int b) {  
    int c;  
    g(a, b, c);  
    return a + b + c;  
}
```

```
void g(int a, int b, int &c) {  
    c = a * b;  
}
```

```
struct S {  
    float x, y, z;  
    double delta;  
  
    double compute();  
};
```

```
double f() {  
    S s;  
    s.x = /* expensive compute */;  
    s.y = /* expensive compute */;  
    s.z = /* expensive compute */;  
    s.delta = s.x - s.y - s.z;  
    return s.compute();  
}
```

Some tips for optimizable APIs:

- Use value semantics! (Already a good idea...)
- Don't create unneeded abstractions. Sometimes, a function parameter is plenty.
- Partition **all** logic away from template-expanded deeply nested constructs.

Use abstractions, but also consider
how they will look to the
optimizer.

Questions!

(maybe answers?)

C++Now
May 2013

Chandler Carruth
chandlerc@gmail.com