

A Discussion of Selected Vienna-Libraries for Computational Science

Karl Rupp¹, Florian Rudolf², Josef Weinbub²



¹MCS Division, Argonne National Laboratory, USA

²Institut für Mikroelektronik, TU Wien, Austria



Pre-Conference Handout Slides (Preview)




C++Now, May 15th, 2013

Part 0: What is this all about?


Simulation Flow

The Many Steps in Simulating a FinFET



Part 1: ViennaCL

Simulation Flow



From Boost.uBLAS to ViennaCL

Consider Existing CPU Code (Boost.uBLAS)

```
using namespace boost::numeric::ublas;

matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

double val = inner_prod(x, y);
y += 2.0 * x;
A += val * outer_prod(x, y);

x = solve(A, y, upper_tag()); // Upper tri. solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```

High-level code with syntactic sugar

From Boost.uBLAS to ViennaCL

Previous Code Snippet Rewritten with ViennaCL

```
using namespace viennacl;
using namespace viennacl::linalg;

matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

double val = inner_prod(x, y);
y += 2.0 * x;
A += val * outer_prod(x, y);

x = solve(A, y, upper_tag()); // Upper tri. solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```

High-level code with syntactic sugar

From Boost.uBLAS to ViennaCL

ViennaCL in Addition Provides Iterative Solvers

```
using namespace viennacl;  
using namespace viennacl::linalg;  
  
compressed_matrix<double> A(1000, 1000);  
vector<double> x(1000), y(1000);  
  
/* Fill A, x, y here */  
  
x = solve(A, y, cg_tag());           // Conjugate Gradients  
x = solve(A, y, bicgstab_tag());      // BiCGStab solver  
x = solve(A, y, gmres_tag());         // GMRES solver
```

No Iterative Solvers Available in Boost.uBLAS...

From Boost.uBLAS to ViennaCL

Thanks to Interface Compatibility

```
using namespace boost::numeric::ublas;
using namespace viennacl::linalg;

compressed_matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

x = solve(A, y, cg_tag());           // Conjugate Gradients
x = solve(A, y, bicgstab_tag());     // BiCGStab solver
x = solve(A, y, gmres_tag());        // GMRES solver
```

Code Reuse Beyond GPU Borders

Eigen <http://eigen.tuxfamily.org/>

MTL 4 <http://www.mtl4.org/>

From Boost.uBLAS to ViennaCL

Generic CG Implementation (Sketch)

```
for (unsigned int i = 0; i < tag.max_iterations(); ++i)
{
    tmp = viennacl::linalg::prod(matrix, p);

    alpha      = ip_rr / inner_prod(tmp, p);
    result     += alpha * p;
    residual   -= alpha * tmp;

    new_ip_rr = inner_prod(residual, residual);
    if (new_ip_rr / norm_rhs_squared < tag.tolerance())
        break;


    beta  = new_ip_rr / ip_rr;
    ip_rr = new_ip_rr;

    p = residual + beta * p;
}
```

About ViennaCL

About

High-level linear algebra C++ library
OpenMP, OpenCL, and CUDA backends
Header-only
Multi-platform



Dissemination

Free Open-Source MIT (X11) License
<http://viennacl.sourceforge.net/>
50-100 downloads per week

Design Rules

Reasonable default values
Compatible to Boost.uBLAS whenever possible
In doubt: clean design over performance

Basic Types

scalar

vector

matrix, compressed_matrix, coordinate_matrix, ell_matrix, hyb_matrix

Data Initialization

Using `viennacl::copy()`

```
std::vector<double>      std_x(100);
ublas::vector<double>    ublas_x(100);
viennacl::vector<double> vcl_x(100);

for (size_t i=0; i<100; ++i) {
    std_x[i] = rand();
    ublas_x[i] = rand();
    vcl_x[i] = rand();    //possible, inefficient
}
```

Basic Types

scalar

vector

matrix, compressed_matrix, coordinate_matrix, ell_matrix, hyb_matrix

Data Initialization

Using `viennacl::copy()`

```
std::vector<double>      std_x(100);
ublas::vector<double>    ublas_x(100);
viennacl::vector<double> vcl_x(100);

/* setup of std_x and ublas_x omitted */

viennacl::copy(std_x.begin(), std_x.end(),
               vcl_x.begin()); //to GPU
viennacl::copy(vcl_x.begin(), vcl_x.end(),
               ublas_x.begin()); //to CPU
```

Basic Types

scalar

vector

matrix, compressed_matrix, coordinate_matrix, ell_matrix, hyb_matrix

Data Initialization

Using `viennacl::copy()`

```
std::vector<std::vector<double> >    std_A;  
ublas::matrix<double>                ublas_A;  
viennacl::matrix<double>              vcl_A;  
  
/* setup of std_A and ublas_A omitted */  
  
viennacl::copy(std_A, vcl_A);          // CPU to GPU  
viennacl::copy(vcl_A, ublas_A);       // GPU to CPU
```

Vector Addition

```
x = y + z;
```

Temporaries are costly (particularly on GPUs)

Expression Templates

Limited expansion

Map to a set of predefined kernels

```
vector_expression<vector<T>, op_plus, vector<T> >  
operator+(vector<T> & v, vector<T> & w) { ... }  
  
vector::operator=(vector_expression<...> const & e) {  
    viennacl::linalg::avbv(*this, 1, e.lhs(), 1, e.rhs());  
}
```


Vector Addition

```
// x = y + z
void avbv(...) {
    switch (active_handle_id(x)) {
        case MAIN_MEMORY:
            host_based::avbv(...);
            break;
        case OPENCL_MEMORY:
            opencl::avbv(...);
            break;
        case CUDA_MEMORY:
            cuda::avbv(...);
            break;
        default:
            raise_error();
    }
}
```

Memory buffers can switch memory domain at runtime

Memory Buffer Migration

```
vector<double> x = zero_vector<double>(42);  
  
memory_types src_memory_loc = memory_domain(x);  
switch_memory_domain(x, MAIN_MEMORY);  
  
/* work on x in main memory here */  
  
switch_memory_domain(x, src_memory_loc);
```




Generalizing compute kernels


```
// x[4:2:8] = a * y[2:2:6] + b * z[3:2:7]
__kernel void avbv(
    double * x, uint off_x, uint inc_x,
    double a,
    double * y, uint off_y, uint inc_y,
    double b,
    double * z, uint off_z, uint inc_z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[off_x + i * inc_x] = a * y[off_y + i * inc_y]
                               + b * z[off_z + i * inc_z];
}
```

No penalty on GPUs because FLOPs are for free

Benchmarks



Benchmarks




Benchmarks

Matrix-Matrix Multiplication

Autotuning environment

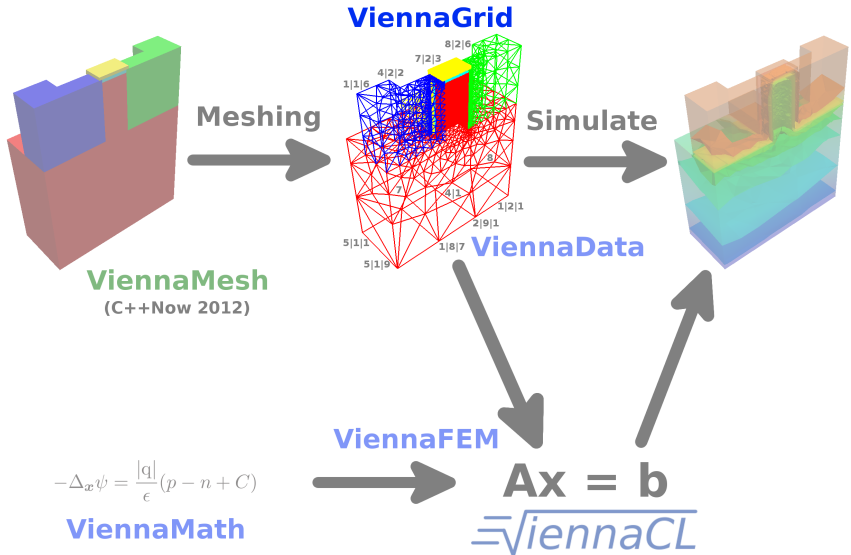
GFLOP Performance for GEMM (Higher is Better)



(AMD Radeon HD 7970, single precision)

Part 2: ViennaGrid

Simulation Flow



What about Boost.Geometry?

A Quick Look at Boost.Geometry

Concepts: Point, Segment, Linestring, Box, etc.

Algorithms: `distance()`, `intersects()`, `convex_hull()`, etc.

Central entity: Point

```
int a[3] = {1, 2, 3};  
int b[3] = {2, 3, 4};  
  
double d = boost::geometry::distance(a, b);
```

Why Boost.Geometry is Not Enough

What about 3D objects? Tetrahedra? Hexahedra?



Traversal of boundary objects missing

No storage facilities for *many* objects

n -cell Concept

Concept of n -cell

Sub- k -cells of an n -cell



Separation of Geometry and Topology

Geometry: Euclidian space, coordinate system



Topology: Connection between points (lines, triangles, etc.)

ViennaGrid Datastructure

Datastructure Requirements

Don't store boundary k -cells unnecessarily

Fast local iteration on k -cells



	Amount	Mem/Obj.	Total Mem.
Vertices	4913	24 B	115 KB
Edges	31024	16 B	485 KB
Facets	50688	48 B	2376 KB
Cells	24576	112 B	2688 KB
Total			5664 KB


	Amount	Mem/Obj.	Total Mem.
Vertices	4913	24 B	115 KB
Edges	0	-	0 KB
Facets	0	-	0 KB
Cells	24576	32 B	768 KB
Total			883 KB

n-cell Implementation

Implementation of `element_t`

Recursive inheritance from boundary layer of dimension $n - 1$

Tag dispatching to enable/disable topological layers



```
template <typename ConfigType, typename ElementTag>
class element_t :
    public boundary_ncell_layer<ConfigType, ElementTag, ElementTag::dim-1>,
    public result_of::element_id_handler<ConfigType, ElementTag>::type
```

User API Design Goals

STL-style, reuse conventions

Allow index-based traversal

Avoid common C++ pitfalls (e.g. template member functions)

```
//iteration over all vertices in the domain:
typedef result_of::ncell_range<DomainType, 0>::type   VertexRange;
typedef result_of::iterator<VertexRange>::type        VertexIterator;

VertexRange vertices = ncells<0>(domain);
for (VertexIterator it = vertices.begin();
     it != vertices.end();
     ++it)
{
    // do something with each vertex here
}
```

Ranges provide iterators over *n*-cells

Extendible

User API Design Goals

Boundary iteration: $k < n$

Coboundary iteration: $k > n$

```
//iteration over all triangles of a vertex
typedef result_of::ncell_range<VertexType, 2>::type TriangleRange;
typedef result_of::iterator<TriangleRange>::type TriangleIterator;

TriangleRange triangles = ncells<0>(vertex, domain);
for (TriangleIterator it = triangles.begin();
      it != triangles.end();
      ++it)
{
    // do something with each triangle here
}
```

Coboundary information not a-priori available in datastructure

Built and cached at first request

Domain Configuration

Top Level Configuration of Domain

Mostly a collection of tags

Predefined configuration classes for common cases

ViennaGrid 1.0.x: Supplemented by global customizations

```
struct triangular_3d
{
    typedef double           numeric_type;
    typedef cartesian_cs<3>   coordinate_system_tag;
    typedef triangle_tag      cell_tag;
};

result_of::domain<triangular_3d>::type      Domain;
```

Work in progress: Everything in a single config class

ViennaGrid Features

Other Features

Segments


I/O: VTK, various mesh-formats

Voronoi information

Refinement


Work in Progress

PLC, hybrid meshes, multigrid



Part 3: ViennaData

Simulation Flow



Plain Object-Oriented Approach?

```
struct Triangle {  
    PointType a, b, c;  
    bool is_on_boudary;  
    double rho; };    //e.g. specific mass
```

Pros and Cons

Data is directly stored with the object

Each (!) triangle carries a boolean flag

Reusability reduced \Rightarrow better rename to `TriangleWithMaterial`

Monolithic, don't do this!

Store Data with Mesh:

```
class Mesh {  
    vector<Triangle> triangles;  
    std::map<size_t, bool> boudary_map; //no memory wasted  
    vector<double> rho_for_triangles;    };
```

Pros and Cons

Triangle is still reusable

Mesh class has to handle data storage complexity

Each additional data requires a change of Mesh class

Mesh object has to be passed to all modules

Monolithic, don't do this!

Approach by ViennaData

Introduction of a hidden data container

Data is stored in a map-like manner using keys

```
//generic interface:  
viennadata::access<KeyType, ValueType>(key) (obj);  
  
//boundary flag and specific mass retrieval for triangle:  
bool on_boundary = access<BoundaryKeyType, bool>(boundary_key) (triangle);  
double rho = access<MassKeyType, double>( mass_key) (triangle);
```

Pros and Cons

Triangle and Mesh are still reusable

Can be used with arbitrary objects (third-party libraries)

Unified interface for data access

Generic Interface

```
viennadata::access<KeyType, ValueType>(key) (obj);
```

Default Storage Scheme

If nothing is known about the object:

```
std::map<ObjectType *,  
         std::map<KeyType, ValueType> >
```

The compiler creates such a map for each of the following:

```
access<long, double>(42) (triangle);  
access<char, double>('b') (triangle);  
access<std::string, double>("mass") (triangle);  
access<std::string, double>("mass") (vertex);
```

Performance Considerations

$\mathcal{O}(\log N) + \mathcal{O}(\log K)$ access time

N ... objects of same type (e.g. triangles)

K ... keys of same type

Usually too slow in a high-performance setting

Can We Do Better?

In general: No

In certain situations: Yes

Type-based Key Dispatch

Only one key per type

$\mathcal{O}(\log N)$ access time

```
access<mass_key, double>(mass_key())(triangle);
```

```
// or:
```

```
access<mass_key, double>()(triangle);
```

Internal Datastructure

```
std::map<ObjectType *, ValueType>
```

Transparent to user

Simple preprocessor macro for activation

Numeric IDs for Objects

Only one key per type

$\mathcal{O}(K)$ access time

$\mathcal{O}(1)$ access time with type-based key dispatch

```
access<std::string, double>("mass") (triangle);
```

```
// with type-based key dispatch:
```

```
access<mass_key, double>() (triangle);
```

Internal Datastructure

```
std::vector<ValueType>
```

Transparent to user

Simple preprocessor macro for activation

Overload generic `id()` accessor

Benchmarking ViennaData


ID-based access to data via ViennaData (class `LightWeight`)
OOP-style storage in classes with payload

```
template <size_t N>
struct FatClass {
    double data;
    char payload[N];
};
```

	10 ³ Objects (us)	10 ⁶ Objects (ms)
LightWeight	4	5
FatClass<10>	1.3	4
FatClass<100>	2.1	11
FatClass<1000>	2.5	11

Part 4: ViennaMath

Simulation Flow



A Symbolic Math Library in C++

Symbolic evaluation and manipulation of math expressions

Unified run time and compile time interface

Differentiation and integration provided

L^AT_EX output

Example Usage

```
variable x(0);  
variable y(1);  
variable z(2);  
expr f = x + y - z;  
expr g = f * f;  
eval(f, make_vector(1.0, 2.0, 4.0)); // returns -1.0  
eval(g, make_vector(1.0, 2.0, 4.0)); // returns 4.0
```

Run Time Evaluation

A Symbolic Math Library in C++

Symbolic evaluation and manipulation of math expressions

Unified run time and compile time interface

Differentiation and integration provided

L^AT_EX output

Example Usage

```
ct_variable<0> x;  
ct_variable<1> y;  
ct_variable<z> z;  
ct_constant<1> c1;  
ct_constant<2> c2;  
ct_constant<4> c4;  
eval(x + y - z, make_vector(c1, c2, c4)); // returns -1
```

Compile Time Evaluation

Substitution and Differentiation (Run Time)

```
variable x(0);  
variable y(1);  
variable z(2);  
substitute(x, y, x*y + z); // returns y*y+z  
diff(x*y + z, x);          // returns y
```

Substitution and Differentiation (Compile Time)

```
ct_variable<0> x;  
ct_variable<1> y;  
ct_variable<2> z;  
substitute(x, y, x*y + z); // returns y*y+z  
diff(x*y + z, x);          // returns y
```

Numerical Integration (Run Time): $\int_0^1 x^2 dx$

```
expr f = integral( make_interval(0, 1), x*x, x );

numerical_quadrature integrator(new gauss_quad_1());
integrator(f);                                // method 1
integrator(make_interval(0, 1), x*x, x);      // method 2
```

Analytical Integration (Compile Time): $\int_0^1 x^2 dx, \int_0^1 \int_0^{1-x} xy dx dy$

```
integrate(make_interval(c0, c1),
          x*y,
          x );    //returns y/2.0
integrate(make_interval(c0, c1),
          integrate( make_interval(c0, c1 - x), x*y, y),
          x);     //returns 1.0/24.0
```

Function Symbols

Represent a function (not evaluable)

Differential Operators

Gradient, Divergence


Symbolic Integration Domain

Specify actual integration domain and integration variables *later*

```
function_symbol u;  
equation eq( laplace(u), 1.0 ); // Poisson equation  
  
function_symbol v;  
expr w = integral(symbolic_interval(),  
                   grad(u) * grad(v));
```


Part 5: ViennaFEM

Simulation Flow



Simulation Flow



Library-Centric Design

ViennaCL for linear solver

ViennaData for data storage

ViennaGrid for mesh handling

ViennaMath for symbolic math

Addresses Come&Go in Academia

Focus on one package

No in-depth knowledge of all packages required

Additional emphasis on good interfaces

ViennaGrid deals with the Mesh

```
DomainType my_domain;  
viennagrid::io::netgen_reader my_reader;  
my_reader(my_domain, "mystructure.mesh");
```

Equation specification via ViennaMath: $\Delta u = -1$

```
equation poisson_eq = viennamath::make_equation(laplace(u), -1);
```

Assembly via ViennaFEM

```
viennafem::pde_assembler fem_assembler;  
fem_assembler(viennafem::make_linear_pde_system(poisson_eq, u),  
              my_domain,  
              system_matrix, load_vector);
```

Linear solver provided by ViennaCL

```
VectorType pde_result  
= viennacl::linalg::solve(system_matrix, load_vector, cg_tag());
```

Lame equation for Linear Elasticity

$$\int_{\Omega} \varepsilon(u) : \sigma(v) \, d\Omega = \int_{\Omega} F \cdot v \, d\Omega \quad \forall v \in \mathcal{V}$$

With $F = (0, 0, 1)^T$:

```
std::vector< Expression > strain = strain_tensor(u);
std::vector< Expression > stress = stress_tensor(v);

Equation weak_form_lame = make_equation(
    integral(symbolic_interval(),
              tensor_reduce( strain, stress )),
    // =
    integral(symbolic_interval(),
              constant(1.0) * v[2])
    );
```

Summary

Software Packages

ViennaCL: GPU-accelerated Linear Algebra

ViennaData: Generic Data Storage

ViennaFEM: Modular High-Level Finite Element Package

ViennaGrid: Generic Mesh Datastructure

ViennaMath: Symbolic Math Kernel

`http://vienna{cl,data,fem,grid,math}@sourceforge.net`

Design Philosophy

Orthogonal Software Design

Convenient High-Level User API

Avoid Dependencies for Fast Migration to other Machines

Free Open Source