

A Discussion of Selected Vienna-Libraries for Computational Science

Karl Rupp¹, Florian Rudolf², Josef Weinbub²



¹MCS Division, Argonne National Laboratory, USA

²Institut für Mikroelektronik, TU Wien, Austria



C++Now, May 15th, 2013

Part 0: What is this all about?

Contents

Libraries

ViennaCL

ViennaData

ViennaFEM

ViennaGrid

ViennaIPD

ViennaMath

ViennaMesh

Applications

ViennaMOS

ViennaProfiler

ViennaSHE

ViennaWD

ViennaX

Covered In This Talk

- Requirements in Computational Science

- GPU Computing

- Providing High-Level Interfaces

- Avoiding Monolithic Code

- Abstractions: From Math to Code (and back)

Not Covered

- C++11

- Maximize use of Boost

The Computational Scientist

A Strange Animal

Goal is science, not to execute software

Codes seldomly designed for 'large scale' upfront

Scientists receive software training from scientists

Performance vs. portability and maintainability


Run code on clusters (headless, old software stack, ...)

Come and go

Basili *et al.*, Understanding the High-Performance-Computing Community: A Software Engineer's Perspective. IEEE Software, 2008.

Simulation Flow


The Many Steps in Simulating a FinFET



(C) Intel, 2011


Simulation Flow

The Many Steps in Simulating a FinFET



Part 1: ViennaCL

Simulation Flow



From Boost.uBLAS to ViennaCL

Consider Existing CPU Code (Boost.uBLAS)

```
using namespace boost::numeric::ublas;

matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

double val = inner_prod(x, y);
y += 2.0 * x;
A += val * outer_prod(x, y);

x = solve(A, y, upper_tag()); // Upper tri. solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```

High-level code with syntactic sugar

From Boost.uBLAS to ViennaCL

Previous Code Snippet Rewritten with ViennaCL

```
using namespace viennacl;
using namespace viennacl::linalg;

matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

double val = inner_prod(x, y);
y += 2.0 * x;
A += val * outer_prod(x, y);

x = solve(A, y, upper_tag()); // Upper tri. solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```

High-level code with syntactic sugar

From Boost.uBLAS to ViennaCL

ViennaCL in Addition Provides Iterative Solvers

```
using namespace viennacl;  
using namespace viennacl::linalg;  
  
compressed_matrix<double> A(1000, 1000);  
vector<double> x(1000), y(1000);  
  
/* Fill A, x, y here */  
  
x = solve(A, y, cg_tag());           // Conjugate Gradients  
x = solve(A, y, bicgstab_tag());      // BiCGStab solver  
x = solve(A, y, gmres_tag());         // GMRES solver
```

No Iterative Solvers Available in Boost.uBLAS...

From Boost.uBLAS to ViennaCL

Thanks to Interface Compatibility

```
using namespace boost::numeric::ublas;
using namespace viennacl::linalg;

compressed_matrix<double> A(1000, 1000);
vector<double> x(1000), y(1000);

/* Fill A, x, y here */

x = solve(A, y, cg_tag());           // Conjugate Gradients
x = solve(A, y, bicgstab_tag());     // BiCGStab solver
x = solve(A, y, gmres_tag());       // GMRES solver
```

Code Reuse Beyond GPU Borders

Eigen <http://eigen.tuxfamily.org/>

MTL 4 <http://www.mtl4.org/>

From Boost.uBLAS to ViennaCL

Generic CG Implementation (Sketch)

```
for (unsigned int i = 0; i < tag.max_iterations(); ++i)
{
    tmp = viennacl::linalg::prod(A, p);

    alpha      = ip_rr / inner_prod(tmp, p);
    result     += alpha * p;
    residual   -= alpha * tmp;

    new_ip_rr = inner_prod(residual, residual);
    if (new_ip_rr / norm_rhs_squared < tag.tolerance())
        break;


    beta  = new_ip_rr / ip_rr;
    ip_rr = new_ip_rr;

    p = residual + beta * p;
}
```

About ViennaCL

About

High-level linear algebra C++ library
OpenMP, OpenCL, and CUDA backends
Header-only
Multi-platform



Dissemination

Free Open-Source MIT (X11) License
<http://viennacl.sourceforge.net/>
50-100 downloads per week

Design Rules

Reasonable default values
Compatible to Boost.uBLAS whenever possible
In doubt: clean design over performance

Basic Types

scalar, vector

matrix, compressed_matrix, coordinate_matrix, ell_matrix, hyb_matrix

Data Initialization

```
std::vector<double>      std_x(100);  
ublas::vector<double>   ublas_x(100);  
viennacl::vector<double> vcl_x(100);  
  
for (size_t i=0; i<100; ++i)  
    // std_x[i] = rand(); // (1)  
    // ublas_x[i] = rand(); // (2)  
    vcl_x[i] = rand(); // (3)
```

(3) is slowest by orders of magnitude!

Basic Types

scalar, vector

matrix, compressed_matrix, coordinate_matrix, ell_matrix, hyb_matrix

Data Initialization

Using `viennacl::copy()`

```
std::vector<double>      std_x(100);
ublas::vector<double>    ublas_x(100);
viennacl::vector<double> vcl_x(100);

/* setup of std_x and ublas_x omitted */

viennacl::copy(std_x.begin(), std_x.end(),
               vcl_x.begin()); //to GPU
viennacl::copy(vcl_x.begin(), vcl_x.end(),
               ublas_x.begin()); //to CPU
```

About ViennaCL

Basic Types

scalar, vector

matrix, compressed_matrix, coordinate_matrix, ell_matrix, hyb_matrix

Data Initialization

Using `viennacl::copy()`

```
std::vector<std::vector<double> >    std_A;  
ublas::matrix<double>                ublas_A;  
viennacl::matrix<double>              vcl_A;  
  
/* setup of std_A and ublas_A omitted */  
  
viennacl::copy(std_A, vcl_A);          // CPU to GPU  
viennacl::copy(vcl_A, ublas_A);       // GPU to CPU
```

Iterator concept doesn't quite work on accelerators

Vector Addition

```
x = y + z;
```

Temporaries are costly (particularly on GPUs)

Expression Templates

Limited expansion

Map to a set of predefined kernels

```
vector_expression<vector<T>, op_plus, vector<T> >  
operator+(vector<T> & v, vector<T> & w) { ... }  
  
vector::operator=(vector_expression<...> const & e) {  
    viennacl::linalg::avbv(*this, 1,e.lhs(), 1,e.rhs());  
}
```


Vector Addition

```
// x = y + z
void avbv(...) {
    switch (active_handle_id(x)) {
        case MAIN_MEMORY:
            host_based::avbv(...);
            break;
        case OPENCL_MEMORY:
            opencl::avbv(...);
            break;
        case CUDA_MEMORY:
            cuda::avbv(...);
            break;
        default:
            raise_error();
    }
}
```

Memory buffers can switch memory domain at runtime

Memory Buffer Migration

```
vector<double> x = zero_vector<double>(42);  
  
memory_types src_memory_loc = memory_domain(x);  
switch_memory_domain(x, MAIN_MEMORY);  
  
/* work on x in main memory here */  
  
switch_memory_domain(x, src_memory_loc);
```



Generalizing compute kernels

```
// x = y + z
__kernel void avbv(
    double * x,

    double * y,

    double * z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[i] = y[i] + z[i];
}
```

Generalizing compute kernels

```
// x = a * y + b * z
__kernel void avbv(
    double * x,
    double a,
    double * y,
    double b,
    double * z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[i] = a * y[i] + b * z[i];
}
```

Generalizing compute kernels


```
// x[4:8] = a * y[2:6] + b * z[3:7]
__kernel void avbv(
    double * x, uint off_x,
    double a,
    double * y, uint off_y,
    double b,
    double * z, uint off_z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[off_x + i] = a * y[off_y + i] + b * z[off_z + i];
}
```


Generalizing compute kernels


```
// x[4:2:8] = a * y[2:2:6] + b * z[3:2:7]
__kernel void avbv(
    double * x, uint off_x, uint inc_x,
    double a,
    double * y, uint off_y, uint inc_y,
    double b,
    double * z, uint off_z, uint inc_z, uint size)
{
    size_t i = get_global_id(0);
    for (; i < size; i += get_global_size())
        x[off_x + i * inc_x] = a * y[off_y + i * inc_y]
                               + b * z[off_z + i * inc_z];
}
```

No penalty on NVIDIA GPUs because FLOPs are for free

Benchmarks



Benchmarks

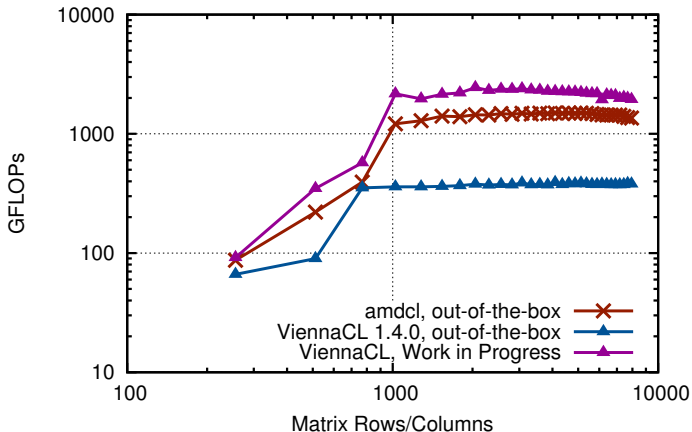


Benchmarks

Matrix-Matrix Multiplication

Autotuning environment

GFLOP Performance for GEMM (Higher is Better)



(AMD Radeon HD 7970, single precision)

Expression Template Limitations


Expression Templates are Not Enough

Consider

```
u = x + y;
```

```
v = x - y;
```

Suboptimal performance with almost any library



Expression Template Limitations


Expression Templates are Not Enough

Consider

```
u = x + y;
```

```
v = x - y;
```

Suboptimal performance with almost any library



Expression Template Limitations

Expression Templates are Not Enough

Consider

```
u = x + y;  
v = x - y;
```

Suboptimal performance with almost any library

OpenCL Kernel Generation


Separate temporary avoidance from operation execution

```
custom_operation op;  
op.add( u = x + y );  
op.add( v = x - y );  
op.execute();           // OpenCL JIT
```

Fully transparent kernel fusion scheduled for release 1.5.0


Expression Template Limitations

Benchmark Results



Part 2: ViennaGrid

Simulation Flow



What about Boost.Geometry?

A Quick Look at Boost.Geometry

Concepts: Point, Segment, Linestring, Box, etc.

Algorithms: `distance()`, `intersects()`, `convex_hull()`, etc.

Central entity: Point

```
int a[3] = {1, 2, 3};  
int b[3] = {2, 3, 4};  
  
double d = boost::geometry::distance(a, b);
```

Why Boost.Geometry is Not Enough

What about 3D objects? Tetrahedra? Hexahedra?



Traversal of boundary objects missing

No storage facilities for *many* objects

n -cell Concept

Concept of n -cell

Sub- k -cells of an n -cell



Separation of Geometry and Topology

Geometry: Euclidian space, coordinate system



Topology: Connection between points (lines, triangles, etc.)

ViennaGrid Datastructure

Datastructure Requirements

Don't store boundary k -cells unnecessarily

Fast local iteration on k -cells



	Amount	Mem/Obj.	Total Mem.
Vertices	4913	24 B	115 KB
Edges	31024	16 B	485 KB
Facets	50688	48 B	2376 KB
Cells	24576	112 B	2688 KB
Total			5664 KB


	Amount	Mem/Obj.	Total Mem.
Vertices	4913	24 B	115 KB
Edges	0	-	0 KB
Facets	0	-	0 KB
Cells	24576	32 B	768 KB
Total			883 KB

n-cell Implementation

Implementation of `element_t`

Recursive inheritance from boundary layer of dimension $n - 1$

Tag dispatching to enable/disable topological layers



```
template <typename ConfigType, typename ElementTag>
class element_t :
    public boundary_ncell_layer<ConfigType, ElementTag, ElementTag::dim-1>,
    public result_of::element_id_handler<ConfigType, ElementTag>::type
```

Domain Configuration

Top Level Configuration of Domain

Mostly a collection of tags

Predefined configuration classes for common cases

ViennaGrid 1.0.x: Supplemented by global customizations

```
struct triangular_3d
{
    typedef double           numeric_type;
    typedef cartesian_cs<3>   coordinate_system_tag;
    typedef triangle_tag      cell_tag;
};

result_of::domain<triangular_3d>::type      Domain;
```

Work in progress: Everything in a single config class

User API Design Goals

STL-style, reuse conventions

Allow index-based traversal

Avoid common C++ pitfalls (e.g. template member functions)

```
//iteration over all vertices in the domain:
typedef result_of::ncell_range<DomainType, 0>::type VertexRange;
typedef result_of::iterator<VertexRange>::type VertexIterator;

VertexRange vertices = ncells<0>(domain);
for (VertexIterator it = vertices.begin();
     it != vertices.end();
     ++it)
{
    // do something with each vertex here
}
```

Ranges provide iterators over n -cells

Extendible

User API Design Goals

STL-style, reuse conventions

Allow index-based traversal

Avoid common C++ pitfalls (e.g. template member functions)

```
//iteration over all vertices in the domain:  
for (std::size_t i = 0; i < ncells<0>(domain).size(); ++i)  
{  
    // do something with ncells<0>(domain)[i] here  
}
```


OpenMP friendly

.size() sometimes known at compiletime!

User API Design Goals

Boundary iteration: $k < n$

Coboundary iteration: $k > n$



```
//iteration over all triangles of a vertex
typedef result_of::ncell_range<VertexType, 2>::type   TriangleRange;
typedef result_of::iterator<TriangleRange>::type      TriangleIterator;

TriangleRange triangles = ncells<2>(vertex, domain);
for (TriangleIterator it = triangles.begin();
      it != triangles.end();
      ++it)
{
    // do something with each triangle here
}
```

Coboundary information not a-priori available in datastructure

Built and cached at first request


ViennaGrid Features

Other Features

- Segments
- I/O: VTK, various mesh-formats
- Voronoi information
- Refinement


Work in Progress

- PLC, hybrid meshes, multigrid



Part 3: ViennaData

Simulation Flow



Plain Object-Oriented Approach?

```
struct Triangle {  
    PointType a, b, c;  
    bool is_on_boudary;  
    double rho; };    //e.g. specific mass
```

Pros and Cons

Data is directly stored with the object

Each (!) triangle carries a boolean flag

Reusability reduced \Rightarrow better rename to `TriangleWithMaterial`

Monolithic, don't do this!

Store Data with Mesh:

```
class Mesh {  
    vector<Triangle> triangles;  
    std::map<size_t, bool> boudary_map; //no memory wasted  
    vector<double> rho_for_triangles;    };
```

Pros and Cons

Triangle is still reusable

Mesh class has to handle data storage complexity

Each additional data requires a change of Mesh class

Mesh object has to be passed to all modules

Monolithic, don't do this!

Approach by ViennaData

Introduction of a hidden data container

Data is stored in a map-like manner using keys

```
//generic interface:  
viennadata::access<KeyType, ValueType>(key) (obj);  
  
//boundary flag and specific mass retrieval for triangle:  
bool on_boundary = access<BoundaryKeyType, bool>(boundary_key) (triangle);  
double rho = access<MassKeyType, double>(mass_key) (triangle);
```

Pros and Cons

- + Triangle and Mesh are still reusable
- + Can be used with arbitrary objects (third-party libraries)
- + Unified interface for data access
- Global state

Generic Interface

```
viennadata::access<KeyType, ValueType>(key) (obj);
```

Default Storage Scheme

If nothing is known about the object:

```
std::map<ObjectType *,  
        std::map<KeyType, ValueType> >
```

The compiler creates such a map for each of the following:

```
access<long, double>(42) (triangle);  
access<char, double>('b') (triangle);  
access<std::string, double>("mass") (triangle);  
access<std::string, double>("mass") (vertex);
```

Performance Considerations

$\mathcal{O}(\log N) + \mathcal{O}(\log K)$ access time

N ... objects of same type (e.g. triangles)

K ... keys of same type

Usually too slow in a high-performance setting

Can We Do Better?

In general: No

In certain situations: Yes

Type-based Key Dispatch

Only one key per type

$\mathcal{O}(\log N)$ access time

```
access<mass_key, double> (mass_key()) (triangle);  
  
// or:  
access<mass_key, double> () (triangle);
```

Internal Datastructure

```
std::map<ObjectType *, ValueType>
```

Transparent to user

One line of code for activation

Numeric IDs for Objects

Only one key per type

$\mathcal{O}(K)$ access time

$\mathcal{O}(1)$ access time with type-based key dispatch

```
access<std::string, double>("mass")(triangle);
```

```
// with type-based key dispatch:
```

```
access<mass_key, double>()(triangle);
```

Internal Datastructure

```
std::vector<ValueType>
```

Transparent to user

One line of code for activation

Overload generic `id()` accessor

Benchmarking ViennaData

ID-based access to data via ViennaData (class `LightWeight`)
OOP-style storage in classes with payload

```
template <size_t N>
struct FatClass {
    double data;
    char payload[N];
};
```

	10 ³ Objects (us)	10 ⁶ Objects (ms)
LightWeight	4	5
FatClass<10>	1.3	4
FatClass<100>	2.1	11
FatClass<1000>	2.5	11

Other Routines

Not further addressed:

```
viennadata::copy<KeyType, ValueType>(key) (src_obj,  
    dst_obj);  
viennadata::move<KeyType, ValueType>(key) (src_obj,  
    dst_obj);  
viennadata::find<KeyType, ValueType>(key) (obj);  
  
viennadata::erase<KeyType, ValueType>(key) (obj);  
viennadata::erase<KeyType, all>(key) (obj);  
viennadata::erase<all, ValueType>(key) (obj);
```

Pitfalls


Inheritance

Limited lifetime of objects

Compilation units with different configuration

Part 4: ViennaMath

Simulation Flow



A Symbolic Math Library in C++

Symbolic evaluation and manipulation of math expressions

Unified run time and compile time interface

Differentiation and integration provided

L^AT_EX output

Example Usage

```
variable x(0);  
variable y(1);  
variable z(2);  
expr f = x + y - z;  
expr g = f * f;  
eval(f, make_vector(1.0, 2.0, 4.0)); // returns -1.0  
eval(g, make_vector(1.0, 2.0, 4.0)); // returns 1.0
```

Run Time Evaluation

A Symbolic Math Library in C++

Symbolic evaluation and manipulation of math expressions

Unified run time and compile time interface

Differentiation and integration provided

L^AT_EX output

Example Usage

```
ct_variable<0> x;  
ct_variable<1> y;  
ct_variable<2> z;  
ct_constant<1> c1;  
ct_constant<2> c2;  
ct_constant<4> c4;  
eval(x + y - z, make_vector(c1, c2, c4)); // returns -1
```

Compile Time Evaluation

Substitution and Differentiation (Run Time)

```
variable x(0);  
variable y(1);  
variable z(2);  
substitute(x, y, x*y + z); // returns y*y+z  
diff(x*y + z, x);          // returns y
```

Substitution and Differentiation (Compile Time)

```
ct_variable<0> x;  
ct_variable<1> y;  
ct_variable<2> z;  
substitute(x, y, x*y + z); // returns y*y+z  
diff(x*y + z, x);          // returns y
```

Numerical Integration (Run Time): $\int_0^1 x^2 dx$

```
expr f = integral( make_interval(0, 1), x*x, x );  
  
numerical_quadrature integrator(new gauss_quad_1());  
integrator(f); // method 1  
integrator(make_interval(0, 1), x*x, x); // method 2
```

Analytical Integration (Compile Time): $\int_0^1 x^2 dx, \int_0^1 \int_0^{1-x} xy dx dy$

```
integrate(make_interval(c0, c1),  
          x*y,  
          x ); //returns y/2.0  
integrate(make_interval(c0, c1),  
          integrate( make_interval(c0, c1 - x), x*y, y),  
          x); //returns 1.0/24.0
```

Function Symbols

Represent a function (not evaluable)

Differential Operators

Gradient, Divergence


Symbolic Integration Domain

Specify actual integration domain and integration variables *later*

```
function_symbol u;  
equation eq( laplace(u), 1.0 ); // Poisson equation  
  
function_symbol v;  
expr w = integral(symbolic_interval(),  
                   grad(u) * grad(v));
```

Part 5: ViennaFEM

Simulation Flow





Library-Centric Design

ViennaCL for linear solver

ViennaData for data storage

ViennaGrid for mesh handling

ViennaMath for symbolic math

Addresses Come&Go in Academia

Focus on one package

No in-depth knowledge of all packages required

Additional emphasis on good interfaces

ViennaGrid deals with the Mesh

```
DomainType my_domain;  
viennagrid::io::netgen_reader my_reader;  
my_reader(my_domain, "mystructure.mesh");
```

Equation specification via ViennaMath: $\Delta u = -1$

```
equation poisson_eq = viennamath::make_equation(laplace(u), -1);
```

Assembly via ViennaFEM

```
viennafem::pde_assembler fem_assembler;  
fem_assembler(viennafem::make_linear_pde_system(poisson_eq, u),  
              my_domain,  
              system_matrix, load_vector);
```

Linear solver provided by ViennaCL

```
VectorType pde_result  
= viennacl::linalg::solve(system_matrix, load_vector, cg_tag());
```

Lame equation for Linear Elasticity


$$\int_{\Omega} \varepsilon(u) : \sigma(v) \, d\Omega = \int_{\Omega} F \cdot v \, d\Omega \quad \forall v \in \mathcal{V}$$

With $F = (0, 0, 1)^T$:

```
std::vector< Expression > strain = strain_tensor(u);
std::vector< Expression > stress = stress_tensor(v);

Equation weak_form_lame = make_equation(
    integral(symbolic_interval(),
              tensor_reduce( strain, stress )),
    // =
    integral(symbolic_interval(),
              constant(1.0) * v[2])
    );
```

Solving Lamé's Equation for a TSV



Summary

Software Packages

- ViennaCL: GPU-accelerated Linear Algebra
- ViennaData: Generic Data Storage
- ViennaFEM: Modular High-Level Finite Element Package
- ViennaGrid: Generic Mesh Datastructure
- ViennaMath: Symbolic Math Kernel

`http://vienna{cl,data,fem,grid,math}@sourceforge.net`

Design Philosophy

- Orthogonal Software Design
- Convenient High-Level User API
- Avoid Unneeded Dependencies
- Free Open Source (MIT License)